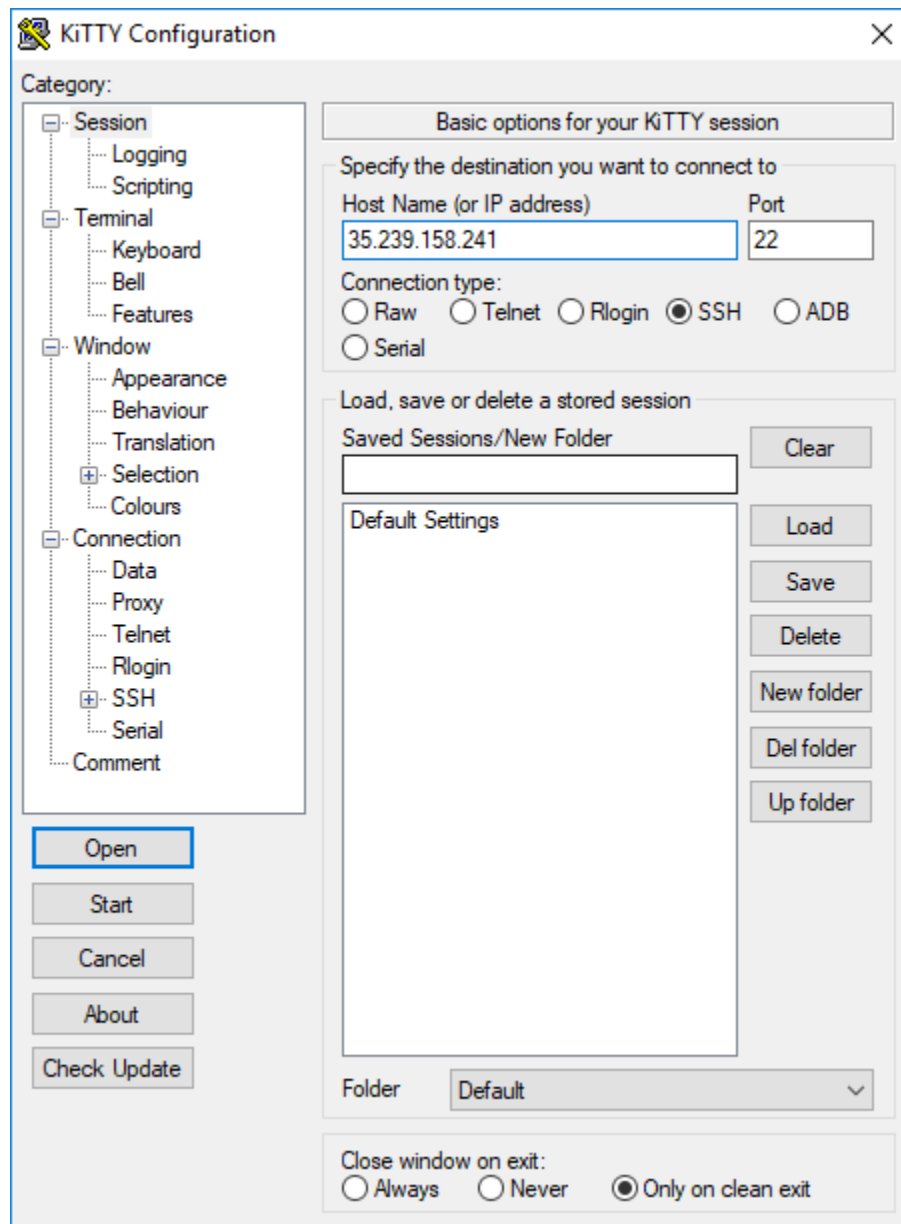


This document was a guideline of sentiment analysis using Spark. We're using the 35.239.158.241:22 ssh through kitty, because of kitty stored all the session if our connection are not stable. Kitty ui was present at the below.



Login with training29 and password Cl0ud3r4* to start the spark from server, use command 'source /tmp/source_profile' to declare the source of spark and type pyspark2 to start the spark.

```
training29@cloudera-master1:~$ source /tmp/source_profile
[training29@cloudera-master1 ~]$ pyspark2
Python 2.7.5 (default, Aug 7 2019, 00:51:29)
[GCC 4.8.5 20150623 (Red Hat 4.8.5-39)] on linux2
Type "help", "copyright", "credits" or "license()" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
19/10/11 02:42:03 WARN util.Utils: Service 'SparkUI' could not bind on port 4040
. Attempting port 4041.
19/10/11 02:42:03 WARN util.Utils: Service 'SparkUI' could not bind on port 4041
. Attempting port 4042.
19/10/11 02:42:03 WARN util.Utils: Service 'SparkUI' could not bind on port 4042
. Attempting port 4043.
Welcome to

  ____      __
 / ___ |__ / /_ __
/ /___/ __/ / / __/
/____/___/_/_/___/
version 2.4.0.cloudera2

Using Python version 2.7.5 (default, Aug 7 2019 00:51:29)
SparkSession available as 'spark'.
>>> █
```

1. Import and Data Preparation

Import the data.

```
>>> df = spark.read.format("csv").option("header", "true").load("/user/cloudera/clean_tweet.csv")
```

```
>>> df.show(5)
+-----+-----+
|          text|target|
+-----+-----+
|awww that s a bum...|    0|
|is upset that he ...|    0|
|i dived many time...|    0|
|my whole body fee...|    0|
|no it s not behav...|    0|
+-----+-----+
only showing top 5 rows
```

Get the columns, and store into cols variable, there was a columns named as text and target. The text column stored a tweet data and the target stored a sentiment, 1 means positive, 0 means negative.

```
>>> col = df.columns
>>> col
['text', 'target']
```

Count the value of target data

```
>>> df.groupBy('target').count().show()
+-----+-----+
|target| count|
+-----+-----+
|      0|800000|
|      1|800000|
+-----+-----+
```

Checking the null value from the data.

```
>>> df.filter(df.target.isNull()).count()
0
>>> df.filter(df.text.isNull()).count()
3247
```

There's null value in the text data, so we must drop that data to improve our classification.

```
>>> df = df.na.drop(subset=['text'])
```

And then, we must check the count of data, to make sure that we done that we do.

```
>>> df.groupby('target').count().show()
+-----+-----+
|target| count|
+-----+-----+
|      0|798503|
|      1|798250|
+-----+-----+
```

2. Feature Extraction

After data cleansing, we gonna be make a token with tokenizer library, and then extract the feature with hashing convectorizer and idf vectorizer.

```
train_df.show(5)>>> from pyspark.ml.feature import StringIndexer
>>> from pyspark.ml import Pipeline
>>>
>>> tokenizer = Tokenizer(inputCol="text", outputCol="words")
>>> hashtf = HashingTF(numFeatures=2**16, inputCol="words", outputCol='tf')
>>> idf = IDF(inputCol='tf', outputCol="features", minDocFreq=5) #minDocFreq: re
move sparse terms
>>> label_stringIdx = StringIndexer(inputCol = "target", outputCol = "label")
>>> pipeline = Pipeline(stages=[tokenizer, hashtf, idf, label_stringIdx])
>>>
>>> pipelineFit = pipeline.fit(train_set)
>>> train_df = pipelineFit.transform(train_set)
>>> val_df = pipelineFit.transform(val_set)
>>> train_df.show(5)
+-----+-----+-----+-----+-----+
+-----+-----+
|          text|target|          words|          tf|
|  features|label|
+-----+-----+-----+-----+-----+
+-----+-----+
|          a|    0|          [a]| (65536, [30802], [1...| (65536, [3
0802], [1...| 0.0|
|a actually don t ...|    0|[a, actually, don...| (65536, [1903,1588...| (65536, [1
903,1588...| 0.0|
|a actually due to...|    0|[a, actually, due...| (65536, [338,1903,...| (65536, [3
38,1903,...| 0.0|
|a ah kan fb nih n...|    0|[a, ah, kan, fb, ...| (65536, [546,6387,...| (65536, [5
46,6387,...| 0.0|
|  a airfranceflight|    0|[a, airfranceflight]| (65536, [30802,527...| (65536, [3
0802,527...| 0.0|
+-----+-----+-----+-----+-----+
+-----+-----+
only showing top 5 rows
```

3. Classification and evaluation

After the feature extracted, we import the machine learning library from `pyspark.ml.classification`, in this implementation we use logistic regression. From the result we get, the accuracy from model was 86 % and the accuracy from data test was 79%.

```
>>> from pyspark.ml.classification import LogisticRegression
>>> lr = LogisticRegression(maxIter=100)
lrModel = lr.fit(train_df)
predictions = lrModel.transform(val_df)
from pyspark.ml.evaluation import BinaryClassificationEvaluator
evaluator = BinaryClassificationEvaluator(rawPredictionCol="rawPrediction")
evaluator.evaluate(predictions)>>> lrModel = lr.fit(train_df)
19/10/12 01:18:26 WARN netlib.BLAS: Failed to load implementation from: com.gith
ub.fommil.netlib.NativeSystemBLAS
19/10/12 01:18:26 WARN netlib.BLAS: Failed to load implementation from: com.gith
ub.fommil.netlib.NativeRefBLAS
>>> predictions = lrModel.transform(val_df)
>>> from pyspark.ml.evaluation import BinaryClassificationEvaluator
>>> evaluator = BinaryClassificationEvaluator(rawPredictionCol="rawPrediction")
>>> evaluator.evaluate(predictions)
0.8613668961813948
>>> accuracy = predictions.filter(predictions.label == predictions.prediction).c
ount() / float(val_set.count())
>>> accuracy
0.7907875562313882
```